

Python Code

Output

Exercise 1:

Impute missing values (categorical variables)

train.Gender.fillna("Female",inplace=True)

train.Married.fillna("No",inplace=True)

Exercise 1:

```
In [8]: train.Gender.fillna("Female",inplace=True)
...: train.Married.fillna("No",inplace=True)

In [9]: print(train.isnull().sum())
Loan_ID      0
Gender        0
Married       0
Dependents    15
Education     0
Self_Employed 32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
Loan_Status   0
dtype: int64
```

Exercise 2:

Impute missing values (numerical variables)

train.fillna(train.median(),inplace=True)

Exercise 2:

```
In [7]: print(train.isnull().sum())
Loan_ID      0
Gender        13
Married       3
Dependents    15
Education     0
Self_Employed 32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
Loan_Status   0
dtype: int64
```

train.isnull().sum()

```
ApplicantIncome    3812.5
CoapplicantIncome   1188.5
LoanAmount         128.0
Loan_Amount_Term    360.0
Credit_History      1.0
```

train.median()

Python Code

Exercise 3:

```
#Split train data for cross
validation
from sklearn.model_selection import
train_test_split
x_train,x_cv,y_train,y_cv =
train_test_split(X,y,test_size=0.3)
```

Exercise 4:

```
#Predict values using kNN
pred_test=kNN.predict(test)

#Write test results in csv file

predictions=pd.DataFrame(pred_test,columns=['pre
dictions']).to_csv('knn_predicton.csv')
```

```
np.unique(pred_test, return_counts=True)
```

Output

Exercise 3:

| | | | |
|---------|-----------|-----------|--|
| x_cv | DataFrame | (185, 20) | Column names: ApplicantIncome, CoapplicantIncome, LoanAmount, Loan_Amo ... |
| x_train | DataFrame | (429, 20) | Column names: ApplicantIncome, CoapplicantIncome, LoanAmount, Loan_Amo ... |
| y | Series | (614,) | Series object of pandas.core.series module |
| y_cv | Series | (185,) | Series object of pandas.core.series module |
| y_train | Series | (429,) | Series object of pandas.core.series module |

Exercise 4:

| | A | B | C | D | E | F |
|----|---|-------------|---------------|-----|---|---|
| 1 | | predictions | | | | |
| 2 | | 0 Y | Number of "Y" | 307 | | |
| 3 | | 1 Y | Number of "N" | 60 | | |
| 4 | | 2 N | | | | |
| 5 | | 3 Y | | | | |
| 6 | | 4 Y | | | | |
| 7 | | 5 Y | | | | |
| 8 | | 6 N | | | | |
| 9 | | 7 Y | | | | |
| 10 | | 8 Y | | | | |
| 11 | | 9 Y | | | | |
| 12 | | 10 Y | | | | |
| 13 | | 11 Y | | | | |
| 14 | | 12 Y | | | | |

```
(array(['N', 'Y'], dtype=object), array([ 60, 307],
dtype=int64))
```

Python Code

Exercise 5:

```
import pandas as pd
import numpy as np

#Load data files
train=pd.read_csv("train.csv")
test=pd.read_csv("test.csv")

#Find missing values
train.isnull().sum()
test.isnull().sum()

#Impute missing values with mean (numerical variables)
train.fillna(train.mean(),inplace=True)
train.isnull().sum()

#Test data
test.fillna(test.mean(),inplace=True)
test.isnull().sum()

#Impute missing values with mode (categorical variables)
train.Gender.fillna(train.Gender.mode()[0],inplace=True)
train.Married.fillna(train.Married.mode()[0],inplace=True)
train.Dependents.fillna(train.Dependents.mode()[0],inplace=True)
train.Self_Employed.fillna(train.Self_Employed.mode()[0],inplace=True)
train.isnull().sum()

#Test data
test.Gender.fillna(test.Gender.mode()[0],inplace=True)
test.Dependents.fillna(test.Dependents.mode()[0],inplace=True)
test.Self_Employed.fillna(test.Self_Employed.mode()[0],inplace=True)
test.isnull().sum()

#Treatment of outliers
train.Loan_Amount_Term=np.log(train.Loan_Amount_Term)

#(3)PREDICTIVE MODELLING
#Remove Loan_ID variable - Irrelevant
train=train.drop("Loan_ID",axis=1)
test=test.drop("Loan_ID",axis=1)

#Create target variable
X=train.drop("Loan_Status",1)
y=train["Loan_Status"]
#Build dummy variables for categorical variables
X=pd.get_dummies(X)
train=pd.get_dummies(train)
test=pd.get_dummies(test)

#Split train data for cross validation
from sklearn.model_selection import train_test_split
x_train,x_cv,y_train,y_cv = train_test_split(X,y,test_size=0.2)

#(a)LOGISTIC REGRESSION ALGORITHM
#Fit model
# from sklearn.linear_model import LogisticRegression
# model=LogisticRegression()
# model.fit(x_train,y_train)

# Random Forest
#Fit model
from sklearn.ensemble import RandomForestClassifier
# Invoke the model
rf=RandomForestClassifier(random_state=42)

# train the model
rf.fit(x_train,y_train)

from sklearn.metrics import roc_auc_score
roc_rf=roc_auc_score(y_cv,rf.predict_proba(x_cv)[0,1])
print("the AUC-ROC score for Random Forest is", roc_rf)

# plot ROC curve
from sklearn.metrics import RocCurveDisplay
import matplotlib.pyplot as plt

logit_disp = RocCurveDisplay.from_estimator(rf, x_cv, y_cv)
ax = plt.gca()
plt.show()
```

Output

Exercise 5:

the AUC-ROC score for Random Forest is
0.7466867469879518

